

Environment Modules

Sebastian Block

26. Oktober 2018

Environment modules are used to provide the dynamic modification of a user's environment via modulefiles.

Easy way to alter or change environment variables such as \$PATH or \$LD_LIBRARY_LOAD.

Be careful not to confuse “environment modules” with “kernel modules.”

module avail

shows the available modules in \$MODULEPATH. Long output with -l

module add|load *Modulename*

load the module with the modulename (output from module avail).

Possible to load multiple modules.

Modules not in \$MODULEPATH can be loaded with the full or relative path. (But not only filename!)

Using modules

rm|unload *Modulename*

unload the module

module purge

unload all modules

module list

list all loaded modules

module show *Modulename*

Display information about one or more modulefiles

module use *Directory*

Add *Directory* to the first position of \$MODULEPATH. To append use -a.

module unuse *Directory*

Remove *Directory* from \$MODULEPATH

Using modules

Examples

```
[sblock@frontend01 ~] module avail -l
```

```
- Package -----+- Versions -+- Last mod. -----
```

```
/cluster/modulefiles:
```

comp/gcc/7.2.0	2017/12/30	16:39:34
julia/0.6.2	2018/06/06	11:24:58
mpi/openmpi/gcc/3.0.0	2017/12/30	17:16:45
mpi/openmpi/gcc/3.1.2	2018/10/09	12:31:01
mpi/openmpi-java/gcc/3.0.0	2018/06/05	14:35:26
mpi/openmpi-x86_64	2017/08/03	20:28:39
nvidia/cuda/9.0	2018/02/07	13:41:18
nvidia/cuda/9.2.88	2018/05/31	8:50:48
singularity/2.5.2	2018/10/12	12:48:52

Using modules

Examples

```
module load mpi/openmpi/gcc/3.1.2 comp/gcc/7.2.0
```

```
module list
```

Currently Loaded Modulefiles:

1) mpi/openmpi/gcc/3.1.2 2) comp/gcc/7.2.0

```
module purge
```

```
module load /home/users/s/sblock/modules/turbomole_test
```

```
module list
```

Currently Loaded Modulefiles:

1) turbomole_test

Using modules

module load can be included in the shell's initialization file (.bash_profile, .bashrc, ...) or in sbatch scripts.

It can be used in bash scripts but then it is only available in the scripts sub-shell.

After loading a module on frontend the environment is used when running srun, salloc or sbatch.

```
[sblock@frontend01 ~]$  
srun gcc --version  
gcc (GCC) 4.8.5 20150623 (Red Hat 4.8.5-16)  
module load comp/gcc/7.2.0  
srun gcc --version  
gcc (GCC) 7.2.0
```


Using modules

module is a function and can not used with srun.

```
srun module list
```

will fail!

To load a module from your current directory use ./modulename.

```
module load ./modulename
```

module can load environment including aliases and other modulefiles. If someone is interested in writing own module files I can prepare a presentation.

- Fairshare works different now. The calculation is still based on used resources. However this is no longer reset every month but calculated with a decay half life time of 7 days.
- The base amount of fairshare shares was changed to respect the higher priority for accounts/users belonging to eecs (Faculty IV).
- The priority weight was changed in some ways:
PriorityWeightAge = 4000
PriorityWeightFairShare = 20000
PriorityWeightJobSize = 1000
PriorityWeightPartition = 30000

- OpenMPI 3.1.2 is now available on the cluster. It can be used with the modulefile `mpi/openmpi/gcc/3.1.2`
- The containerization program Singularity version 2.5.2 is now available and loadable with the modulefile `singularity/2.5.2`
- The OmniPath hardware on SMP001 was replaced.
- The NUMA (Non-uniform memory access) configuration is now set correct on node070 and node132

An easy way to access the files on the cluster storage is *sshfs*. This way you can mount remote directories over ssh.

- 1 create a local mount point

```
mkdir ~/clustermnt
```

- 2 mount the remote folder:

```
sshfs user@gateway.hpc.tu-berlin.de: ~/clustermnt/
```

- 3 the user cluster home folder is now mounted to ~/clustermnt/

- 4 To unmount close all open files in ~/clustermnt and leave the directory in the shell or your file manager. Then umount:

```
fusermount -u ~/clustermnt
```

The remote folder to mount can be relative to home with `:directory` or the absolute path `:/path/to/directory`.

Editing or creating files inside the sshfs-mount is done as user (on the cluster), not the local user.

sshfs is easy but not fast.

correction of the last presentation

The nodes `gpu[001-020]` have **512GB** of memory, not 256GB.
`gpu21` also has two Nvidia Tesla but only 256GB of memory.